

Linus Dominic Nathaniel

Lead Data Engineer / Data Architect

linusnathaniel@gmail.com | +92 323 213 2371 | Karachi, Pakistan
linkedin.com/in/linus-dominic | github.com/linusdominic | linusdominic.github.io
Open to EU relocation (visa sponsorship required) and to EU-remote roles

SUMMARY

Data engineer and architect with 6+ years building platforms on top of messy operational systems. I currently lead a team of 4–6 engineers at Altair Capital Group, a US real-estate finance firm, where I own the data architecture from source-system contracts through warehouse modeling, orchestration, BI and the integrations that feed all of it. Most of my work comes down to reconciliation: getting several systems that each think they own the customer to agree on one version of them. Day to day that means Redshift, BigQuery and Athena over S3 parquet, Airflow, dbt, Spark and Glue. Before Altair I worked in-house at two product companies, including a multi-vendor marketplace platform.

LEADERSHIP

Team Lead, Altair Capital Group (Dec 2025 – present). 4–6 engineers across data platform, BI and integrations.

- I run architecture and code review, plan the sprints, and own the standards and PR conventions the team builds to.
- Every platform initiative listed under Altair below was delivered by this team with me leading the design and the build.
- Mentoring engineers into owning their own domains end to end. Involved in hiring and technical screening.
- I sit between lending and operations stakeholders and the engineers: gathering requirements, negotiating what each system promises to deliver, and getting people onto self-serve analytics instead of ad-hoc requests.

Earlier: ran the Talend to Airflow migration at VentureDive and mentored the junior engineers there. Lead Support Engineer at Bleed AI, running the data-analysis side of computer-vision delivery.

EXPERIENCE

Lead Data Engineer / Data Architect, Altair Capital Group

December 2025 – Present | Remote, US real-estate finance firm | Leading 4–6 engineers

- Built the company's unified customer-journey model, which reconciles a parquet data lake, dialer logs and the CRM into one phone-keyed six-stage funnel across roughly 2M records. It is now the source everything downstream reads from.
- Rewrote the S3 data lake layout: canonical zones, sane partitioning, proper cataloging across about 29 ingestion and transformation jobs. Cleared out years of duplicated data in the process, and the zoning convention is still what the team follows.
- Built the reconciliation framework between the loan-origination platform and the warehouse, so operations can see exactly where the two disagree. Diffs are auditable across millions of records; before this the gaps were simply invisible.
- Moved BI onto Apache Superset and made it deployable from code, using Python and Jinja-templated MySQL views. It powers the Performance Score Ratings, Loan Officer pipeline and Daily Call dashboards that sales, operations and leadership use daily.
- Designed the voice and CRM ingestion pipeline on AWS Lambda, with an S3 pending queue and OAuth tokens cached in SSM and shared across invocations. The previous version kept tripping the provider's rate limits under concurrency; this one does not. Anomaly alerts go out through CloudWatch and SNS.
- Architected the conversation-intelligence layer behind the company's AI telemarketing agent: an LLM detector for meaningful conversations, plus a call-note generator that writes a different note for the Loan Officer, the Processor and the CCO from the same call, straight back into CRM Leads and Contacts.
- Owned the internal loan-pricing platform for the Bridge and DSCR products, from requirements with the lending team through to term-sheet output, including the data model, the calculation engine and the CRM integration.

Senior Data Engineer, VentureDive

June 2025 – December 2025 | Karachi, Pakistan

- Ran the migration of the team's ETL platform off Talend and onto Airflow, along with the standards that came with it: scheduling conventions, observability, DAGs versioned in Git. Redesigning for parallelism brought average runtime down and, more usefully, cut the paging.
- Made Airflow plus dbt plus warehouse the team's reference architecture, with dbt models, tests, sources and lineage adopted across downstream transformations.
- Wrote the multi-cloud deployment standard on Terraform and GitHub Actions, used across AWS, GCP and Azure engagements (S3/Redshift, BigQuery/Dataflow, Synapse/Data Factory).
- Mentored the junior engineers and turned our Airflow patterns, review habits and DAG-testing into written team practice.

Data Engineer, Pure Harvest Smart Farms (product-based agri-tech company)

June 2023 – May 2025 | Remote

- Wrote the ETL behind the company's own product platform in PySpark on AWS Glue. Processing time came down around 30%.
- Automated ingestion from several sources into one AWS data lake, which widened what operations and the growers could actually analyze.
- Built out a Medallion lakehouse (bronze, silver, gold) on Redshift: modeling, governance and a fair amount of performance tuning.
- Built the Superset dashboards used across cultivation and supply chain.

Data and Machine Learning Engineer, Arpatech (Pvt) Ltd (product company, marketplace platform)

January 2022 – June 2023 | Karachi, Pakistan

- Built the Custom Data Insights platform for Connection Cloud Marketplace, a multi-vendor cloud-services marketplace. It ingested vendor and buyer transaction data and gave both sides of the marketplace spend analytics and side-by-side comparison of offerings.
- Built an Elasticsearch and Kibana app for the sales team. Retrieval got roughly 50% faster.
- Shipped a Smart Mirror IoT product doing facial-recognition attendance, which took the manual error out of tracking.

Lead Support Engineer and Data Analyst, Bleed AI

December 2019 – November 2021 | Karachi, Pakistan

- Ran the data-analysis side of the computer-vision projects and fed findings back into model tuning for better image recognition.
- Cleaned up the data pipelines for speed without giving up integrity.

SKILLS

- **Architecture and governance:** dimensional / Kimball modeling, data contracts, source-system contracts, lineage, MDM patterns, Medallion architecture, data domains, conformed dimensions
- **Languages:** Python, SQL, Bash. **Orchestration and transformation:** Apache Airflow, dbt, Prefect, AWS Glue, AWS Step Functions
- **Processing:** Apache Spark, PySpark, Pandas, Apache Kafka, Hadoop
- **Warehouse and lakehouse:** Amazon Redshift, Google BigQuery, Snowflake, AWS Athena, Databricks
- **AWS:** S3, Glue, Athena, Redshift, Lambda, RDS, Step Functions, EventBridge, ECS, Fargate, Secrets Manager. **GCP:** BigQuery, Dataflow, Cloud Storage. **Azure:** Data Factory, Synapse Analytics
- **Databases:** MySQL, PostgreSQL, MongoDB, Elasticsearch. **Infrastructure:** Terraform, Docker, GitHub Actions, CI/CD
- **BI:** Apache Superset, Power BI, Grafana, Kibana, Streamlit
- **LLM work:** conversation intelligence, note generation, prompt engineering, role-aware agent design
- **Integrations:** Zoho CRM API, REST/JSON ingestion, Selenium scrapers

EDUCATION

BSc Computer Science
Karachi Institute of Economics and Technology
August 2018 – August 2022

CERTIFICATIONS

- Certified Data Engineer, Karachi AI
- Certified Data Analyst, Karachi AI
- Computer Vision and Image Processing, Bleed AI

LANGUAGES

- English, fluent
- Urdu, native